

Working with JSON and Array data in BigQuery 2.5 | Google Cloud Skills Boost

Qwiklabs : 20-25 minutes

Overview

BigQuery is Google's fully managed, NoOps, low cost analytics database. With BigQuery you can query terabytes and terabytes of data without having any infrastructure to manage or needing a database administrator. BigQuery uses SQL and can take advantage of the pay-as-you-go model. BigQuery allows you to focus on analyzing data to find meaningful insights.

This lab is an in-depth walkthrough of working with semi-structured data (ingesting JSON, Array data types) inside of BigQuery. Denormalizing your schema into a single table with nested and repeated fields can yield performance improvements, but the SQL syntax for working with array data can be tricky. You will practice loading, querying, troubleshooting, and unnesting various semi-structured datasets.

Objectives

In this lab, you learn about the following:

- Loading semi-structured JSON into BigQuery
- Creating and querying arrays
- Creating and querying structs
- Querying nested and repeated fields

Set up and requirements

For each lab, you get a new Google Cloud project and set of resources for a fixed time at no cost.

1. Sign in to Qwiklabs using an **incognito window**.
2. Note the lab's access time (for example, 1:15:00), and make sure you can finish within that time. There is no pause feature. You can restart if needed, but you have to start at the beginning.
3. When ready, click **Start lab**.
4. Note your lab credentials (**Username** and **Password**). You will use them to sign in to the Google Cloud Console.
5. Click **Open Google Console**.

6. Click **Use another account** and copy/paste credentials for **this** lab into the prompts.

If you use other credentials, you'll receive errors or **incur charges**.

7. Accept the terms and skip the recovery resource page.

Open BigQuery Console

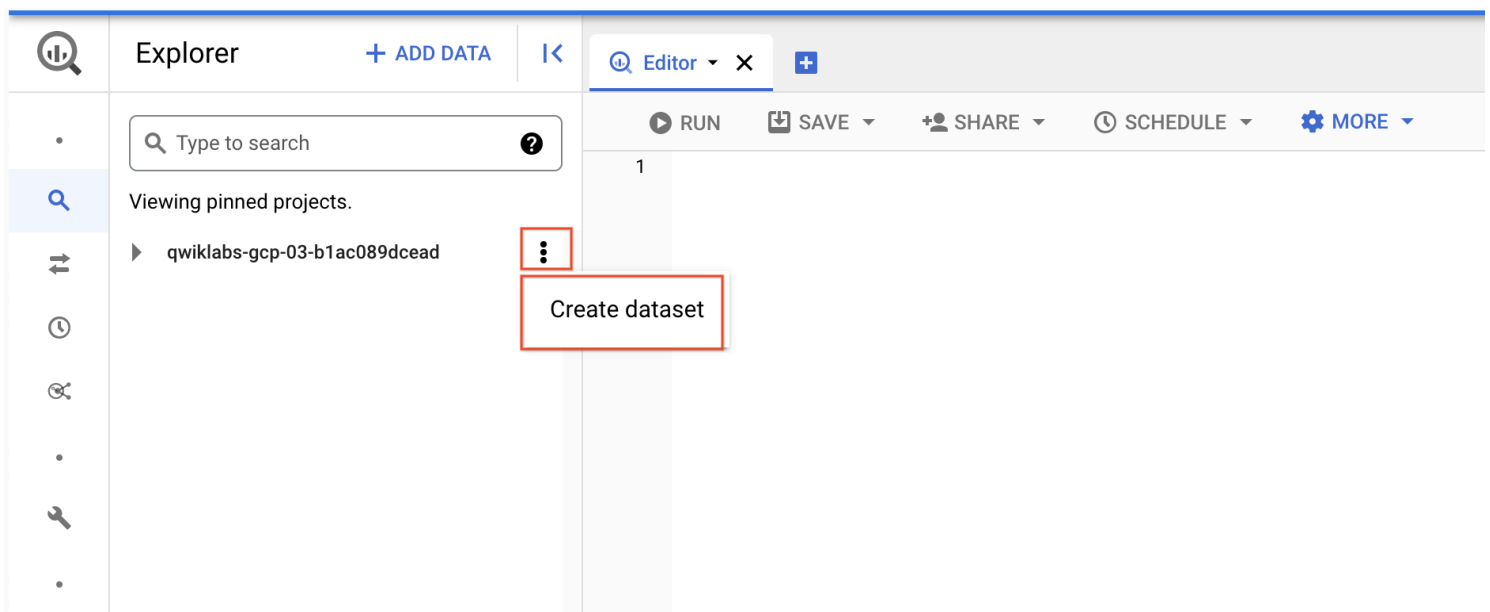
1. In the Google Cloud Console, select **Navigation menu** > **BigQuery**.

The **Welcome to BigQuery in the Cloud Console** message box opens. This message box provides a link to the quickstart guide and lists UI updates.

2. Click **Done**.

Task 1. Create a new dataset to store our tables

1. To create a dataset, click on the **View actions** icon next to your Project ID and then select **Create dataset**.



2. Name the new dataset **fruit_store**. Leave the other options at their default values (Data Location, Default Table Expiration).

3. Click **Create dataset**.

Task 2. Practice working with Arrays in SQL

Normally in SQL you will have a single value for each row like this list of fruits below:

Row Fruit

- 1 raspberry
- 2 blackberry
- 3 strawberry
- 4 cherry

What if you wanted a list of fruit items for each person at the store? It could look something like this:

	Row	Fruit	Person
--	-----	-------	--------

- | | | |
|---|------------|-----------|
| 1 | raspberry | sally |
| 2 | blackberry | sally |
| 3 | strawberry | sally |
| 4 | cherry | sally |
| 5 | orange | frederick |
| 6 | apple | frederick |

In traditional relational database SQL, you would look at the repetition of names and immediately think to split the above table into two separate tables: Fruit Items and People. That process is called [normalization](#) (going from one table to many). This is a common approach for transactional databases like MySQL.

For data warehousing, data analysts often go the reverse direction (denormalization) and bring many separate tables into one large reporting table.

Now, you're going to learn a different approach that stores data at different levels of granularity all in one table using repeated fields:

	Row	Fruit (array)	Person
--	-----	---------------	--------

	raspberry	sally
	blackberry	
1	strawberry	
	cherry	
	orange	frederick
2	apple	

What looks strange about the previous table?

- It's only two rows.
- There are multiple field values for Fruit in a single row.
- The people are associated with all of the field values.

What's the key insight? The array data type!

An easier way to interpret the Fruit array:

Row	Fruit (array)	Person
1	[raspberry, blackberry, strawberry, cherry]	sally
2	[orange, apple]	frederick

Both of these tables are exactly the same. There are two key learnings here:

- An array is simply a list of items in brackets []
- BigQuery visually displays arrays as *flattened*. It simply lists the value in the array vertically (note that all of those values still belong to a single row)

Try it yourself.

1. Enter the following in the BigQuery Query Editor:

```
#standardSQL SELECT ['raspberry', 'blackberry', 'strawberry', 'cherry'] AS fruit_array
```

2. Click **Run**.

3. Now try executing this one:

```
#standardSQL SELECT ['raspberry', 'blackberry', 'strawberry', 'cherry', 1234567] AS fruit_array
```

You should get an error that looks like the following:

```
Error: Array elements of types {INT64, STRING} do not have a common supertype at [3:1]
```

Arrays can only share one data type (all strings, all numbers).

4. Here's the final table to query against:

```
#standardSQL SELECT person, fruit_array, total_cost FROM `data-to-insights.advanced.fruit_store`;
```

5. Click **Run**.

6. After viewing the results, click the **JSON** tab to view the nested structure of the results.

Loading semi-structured JSON into BigQuery

What if you had a JSON file that you needed to ingest into BigQuery?

1. Create a new table in the `fruit_store` dataset.
2. To create a table, click on the **View actions** icon next to the **fruit_store** dataset and select **Create Table**.

Explorer + ADD DATA <

*Unsaved ...ery

RUN

1 #standardS
2 SELECT per

Type to search ?

Viewing pinned projects.

▼ qwiklabs-gcp-03-b1ac089dcead

▶ fruit_store

MORE RESULTS

Open in current tab
Open in new tab
Open in split tab
Delete
Share
Create table

3. Add the following details for the table:

- **Source:** Choose **Google Cloud Storage** in the **Create table from** dropdown.
- **Select file from GCS bucket** (type or paste the following): data-insights-course/labs/optimizing-for-performance/shopping_cart.json
- **File format:** JSONL (Newline delimited JSON) {This will be auto-populated}
- **Schema:** Check **Auto detect** (Schema and input parameters).

4. Call the new table "fruit_details".

5. Click **Create table**.

In the schema, note that `fruit_array` is marked as `REPEATED` which means it's an array.

Recap

- BigQuery natively supports arrays
- Array values must share a data type
- Arrays are called `REPEATED` fields in BigQuery

Click *Check my progress* to verify the objective. Create a new dataset and load JSON data into the table

Task 3. Creating your own arrays with `ARRAY_AGG()`

Don't have arrays in your tables already? You can create them!

1. **Copy and Paste** the below query to explore this public dataset:

```
SELECT fullVisitorId, date, v2ProductName, pageTitle FROM `data-to-insights.ecommerce.all_sessions`  
WHERE visitId = 1501570398 ORDER BY date
```

2. Click **Run** and view the results.

3. Now, we will use the `ARRAY_AGG()` function to aggregate our string values into an array. Copy and paste the below query to explore this public dataset:

```
SELECT fullVisitorId, date, ARRAY_AGG(v2ProductName) AS products_viewed, ARRAY_AGG(pageTitle) AS  
pages_viewed FROM `data-to-insights.ecommerce.all_sessions` WHERE visitId = 1501570398 GROUP BY  
fullVisitorId, date ORDER BY date
```

4. Click **Run** and view the results.

5. Next, we will use the `ARRAY_LENGTH()` function to count the number of pages and products that were viewed:

```
SELECT fullVisitorId, date, ARRAY_AGG(v2ProductName) AS products_viewed,  
ARRAY_LENGTH(ARRAY_AGG(v2ProductName)) AS num_products_viewed, ARRAY_AGG(pageTitle) AS  
pages_viewed, ARRAY_LENGTH(ARRAY_AGG(pageTitle)) AS num_pages_viewed FROM `data-to-  
insights.ecommerce.all_sessions` WHERE visitId = 1501570398 GROUP BY fullVisitorId, date ORDER BY  
date
```

6. Next, let's deduplicate the pages and products so we can see how many unique products were viewed. We'll simply add `DISTINCT` to our `ARRAY_AGG()`:

```
SELECT fullVisitorId, date, ARRAY_AGG(DISTINCT v2ProductName) AS products_viewed,  
ARRAY_LENGTH(ARRAY_AGG(DISTINCT v2ProductName)) AS distinct_products_viewed,  
ARRAY_AGG(DISTINCT pageTitle) AS pages_viewed, ARRAY_LENGTH(ARRAY_AGG(DISTINCT pageTitle))
```

```
AS distinct_pages_viewed FROM `data-to-insights.ecommerce.all_sessions` WHERE visitId = 1501570398
GROUP BY fullVisitorId, date ORDER BY date
```

Recap

You can do some pretty useful things with arrays like:

- Finding the number of elements with `ARRAY_LENGTH(<array>)`
- Deduplicating elements with `ARRAY_AGG(DISTINCT <field>)`
- Ordering elements with `ARRAY_AGG(<field> ORDER BY <field>)`
- Limiting `ARRAY_AGG(<field> LIMIT 5)`

Click *Check my progress* to verify the objective. Creating arrays with `ARRAY_AGG()`

Task 4. Querying datasets that already have ARRAYS

The BigQuery Public Dataset for Google Analytics `bigquery-public-data.google_analytics_sample` has many more fields and rows than our course dataset `data-to-insights.ecommerce.all_sessions`. More importantly, it already stores field values like products, pages, and transactions natively as ARRAYS.

1. **Copy and Paste** the below query to explore the available data and see if you can find fields with repeated values (arrays):

```
SELECT * FROM `bigquery-public-data.google_analytics_sample.ga_sessions_20170801` WHERE visitId = 1501570398
```

2. **Run** the query.

3. **Scroll right** in the results until you see the `hits.product.v2ProductName` field (we will discuss the multiple field aliases shortly).

4. The amount of fields available in the Google Analytics schema can be overwhelming for our analysis. Let's try to query just the visit and page name fields like we did before:

```
SELECT visitId, hits.page.pageTitle FROM `bigquery-public-data.google_analytics_sample.ga_sessions_20170801` WHERE visitId = 1501570398
```

You will get an error: Cannot access field product on a value with type ARRAY> at [5:8]

Before we can query REPEATED fields (arrays) normally, you must first break the arrays back into rows.

For example, the array for `hits.page.pageTitle` is stored currently as a single row like:

```
['homepage','product page','checkout']
```

and we need it to be:

```
['homepage', 'product page', 'checkout']
```


How do we do that with SQL?

5. **Answer:** Use the UNNEST() function on your array field:

```
SELECT DISTINCT visitId, h.page.pageTitle FROM `bigquery-public-  
data.google_analytics_sample.ga_sessions_20170801`, UNNEST(hits) AS h WHERE visitId = 1501570398  
LIMIT 10
```

We'll cover UNNEST() more in detail later but for now just know that:

- You need to UNNEST() arrays to bring the array elements back into rows
- UNNEST() always follows the table name in your FROM clause (think of it conceptually like a pre-joined table)

Click *Check my progress* to verify the objective. Querying datasets that already have ARRAYS

Task 5. Introduction to STRUCTs

You may have wondered why the field alias `hit.page.pageTitle` looks like three fields in one separated by periods. Just as ARRAY values give you the flexibility to *go deep* into the granularity of your fields, another data type allows you to *go wide* in your schema by grouping related fields together. That SQL data type is the **STRUCT** data type.

The easiest way to think about a STRUCT is to consider it conceptually like a separate table that is already pre-joined into your main table.

A STRUCT can have:

- One or many fields in it
- The same or different data types for each field
- Its own alias

Sounds just like a table right?

Let's explore a dataset with STRUCTs

1. Under **Explorer** find the **bigquery-public-data** dataset.
2. If it's not present already, click **ADD DATA > Pin a project**.
3. Click **Enter project name**.
4. Enter **bigquery-public-data** and click **Pin**.
5. Click **bigquery-public-data** in the pinned project list to expand it.
6. Find and open **google_analytics_sample**.
7. Click the **ga_sessions** table.

8. Start scrolling through the schema and answer the following question by using the find feature of your browser (i.e. CTRL + F). **Tip: Expand all columns before you begin counting.**
9. As you can imagine, there is an incredible amount of website session data stored for a modern ecommerce website. The main advantage of having 32 STRUCTs in a single table is it allows you to run queries like this one without having to do any JOINS:

```
SELECT visitId, totals.*, device.* FROM `bigquery-public-data.google_analytics_sample.ga_sessions_20170801` WHERE visitId = 1501570398 LIMIT 10
```

Note: The ``.*`` syntax tells BigQuery to return all fields for that STRUCT (much like it would if `totals.*` was a separate table we joined against).

Storing your large reporting tables as STRUCTs (pre-joined "tables") and ARRAYS (deep granularity) allows you to:

- Gain significant performance advantages by avoiding 32 table JOINS
- Get granular data from ARRAYS when you need it but not be punished if you don't (BigQuery stores each column individually on disk)
- Have all the business context in one table as opposed to worrying about JOIN keys and which tables have the data you need

Click *Check my progress* to verify the objective. Explore a dataset with STRUCTs

Task 6. Practice with STRUCTs and ARRAYS

The next dataset will be lap times of runners around the track. Each lap will be called a "split".



1. With this query, try out the STRUCT syntax and note the different field types within the struct container:

```
#standardSQL SELECT STRUCT("Rudisha" as name, 23.4 as split) as runner
```

Row runner.name runner.split

Row	runner.name	runner.split
1	Rudisha	23.4

What do you notice about the field aliases? Since there are fields nested within the struct (name and split are a subset of runner) you end up with a dot notation.

What if the runner has multiple split times for a single race (like time per lap)?

2. With an array of course! Run the below query to confirm:

```
#standardSQL SELECT STRUCT("Rudisha" as name, [23.4, 26.3, 26.4, 26.1] as splits) AS runner
```

Row runner.name runner.splits

Row	runner.name	runner.splits
1	Rudisha	[23.4, 26.3, 26.4, 26.1]

To recap:

- Structs are containers that can have multiple field names and data types nested inside.
- Arrays can be one of the field types inside of a Struct (as shown above with the splits field).

Practice ingesting JSON data

1. Create a new dataset titled **racing**.
2. Create a new table titled **race_results**.
3. Ingest this Google Cloud Storage JSON file:

[data-insights-course/labs/optimizing-for-performance/race_results.json](#)

- **Source:** Google Cloud Storage under **Create table from** dropdown.

- **Select file from GCS bucket:** data-insights-course/labs/optimizing-for-performance/race_results.json
- **File format:** JSONL (Newline delimited JSON)
- In **Schema**, move the **Edit as text** slider and add the following:

```
[ { "name": "race", "type": "STRING", "mode": "NULLABLE" }, { "name": "participants", "type": "RECORD", "mode": "REPEATED", "fields": [ { "name": "name", "type": "STRING", "mode": "NULLABLE" }, { "name": "splits", "type": "FLOAT", "mode": "REPEATED" } ] } ]
```

4. Click **Create table**.

5. After the load job is successful, preview the schema for the newly created table:

race_results
QUERY
SHARE
COPY
DELETE
LOAD DATA
EXPORT

SCHEMA
DETAILS
PREVIEW

Table schema

Filter Enter property name or value

Field name	Type	Mode	Policy Tags	Description
race	STRING	NULLABLE		
▼ participants	RECORD	REPEATED		
name	STRING	NULLABLE		
splits	FLOAT	REPEATED		

EDIT SCHEMA

Which field is the STRUCT? How do you know?

The **participants** field is the STRUCT because it is of type RECORD.

Which field is the ARRAY?

The participants.splits field is an array of floats inside of the parent participants struct. It has a REPEATED Mode which indicates an array. Values of that array are called nested values since they are multiple values inside of a single field.

Practice querying nested and repeated fields

1. Let's see all of our racers for the 800 Meter race:

```
#standardSQL SELECT * FROM racing.race_results
```

How many rows were returned?

Answer: 1

Query complete (1.236 sec elapsed, 336 B processed)

Job information

Results

JSON

Execution details

Row	race	participants.name	participants.splits
1	800M	Rudisha	23.4
			26.3
			26.4
			26.1
		Makhloufi	24.5
			25.4
			26.6
			26.1
		Murphy	23.9
			26.0
			27.0
			26.0
		Bosse	23.6



What if you wanted to list the name of each runner and the type of race?

2. Run the below schema and see what happens:

```
#standardSQL SELECT race, participants.name FROM racing.race_results
```

Error: Cannot access field name on a value with type ARRAY<STRUCT<name STRING, splits ARRAY<FLOAT64>>>> at [1:21]

Much like forgetting to GROUP BY when you use aggregation functions, here there are two different levels of granularity. One row for the race and three rows for the participants names. So how do you change this:

Row race participants.name

1 800M Rudisha

2 ??? Makhloufi

3 ??? Murphy

...to this:

Row race participants.name

1 800M Rudisha

2 800M Makhloufi

3 800M Murphy

In traditional relational SQL, if you had a races table and a participants table what would you do to get information from both tables? You would JOIN them together. Here the participant STRUCT (which is conceptually very similar to a table) is already part of your races table but is not yet correlated correctly with your non-STRUCT field "race".

Can you think of what two word SQL command you would use to correlate the 800M race with each of the racers in the first table?

Answer: CROSS JOIN

Great!

3. Now try running this:

```
#standardSQL SELECT race, participants.name FROM racing.race_results CROSS JOIN participants # this is the STRUCT (it's like a table within a table)
```

Error: Table name "participants" cannot be resolved: dataset name is missing.

Even though the participants STRUCT is like a table, it is still technically a field in the racing.race_results table.

4. Add the dataset name to the query:

```
#standardSQL SELECT race, participants.name FROM racing.race_results CROSS JOIN racing.race_results.participants # full STRUCT name
```

5. Click **Run**.

Wow! You've successfully listed all of the racers for each race!

Row race name

1 800M Rudisha

2 800M Makhloufi

3 800M Murphy

4 800M Bosse

5 800M Rotich

6 800M Lewandowski

7 800M Kipketer

8 800M Berian

You can simplify the last query by:

- Adding an alias for the original table
- Replacing the words "CROSS JOIN" with a comma (a comma implicitly cross joins)

This will give you the same query result:

```
#standardSQL SELECT race, participants.name FROM racing.race_results AS r, r.participants
```

If you have more than one race type (800M, 100M, 200M), wouldn't a CROSS JOIN just associate every racer name with every possible race like a cartesian product?

Answer: No. This is a *correlated* cross join which only unpacks the elements associated with a single row. For a greater discussion, see [working with ARRAYs and STRUCTs](#).

Click *Check my progress* to verify the objective. Practice with STRUCTs and ARRAYs

Recap of STRUCTs:

- A SQL **STRUCT** is simply a container of other data fields which can be of different data types. The word struct means data structure. Recall the example from earlier:
- `__STRUCT(__"Rudisha" as name, [23.4, 26.3, 26.4, 26.1] as splits__)__AS runner`
- STRUCTs are given an alias (like runner above) and can conceptually be thought of as a table inside of your main table.
- STRUCTs (and ARRAYs) must be unpacked before you can operate over their elements. Wrap an UNNEST() around the name of the struct itself or the struct field that is an array in order to unpack and flatten it.

Task 7. Lab question: STRUCT()

Answer the below questions using the racing.race_results table you created previously.

Task: Write a query to COUNT how many racers were there in total.

- To start, use the below partially written query:

```
#standardSQL SELECT COUNT(participants.name) AS racer_count FROM racing.race_results
```

Note: Remember you will need to cross join in your struct name as an additional data source after the FROM.

Possible solution:

```
#standardSQL SELECT COUNT(p.name) AS racer_count FROM racing.race_results AS r,
UNNEST(r.participants) AS p
```

Answer: There were 8 racers who ran the race.

Task 8. Lab question: Unpacking ARRAYs with UNNEST()

Write a query that will list the total race time for racers whose names begin with R. Order the results with the fastest total time first. Use the UNNEST() operator and start with the partially written query below.

- Complete the query:

```
#standardSQL SELECT p.name, SUM(split_times) as total_race_time FROM racing.race_results AS r ,
r.participants AS p , p.splits AS split_times WHERE GROUP BY ORDER BY ; Note:
```

- You will need to unpack both the struct and the array within the struct as data sources after your FROM clause.
- Be sure to use aliases where appropriate.

Possible solution:

```
#standardSQL SELECT p.name, SUM(split_times) as total_race_time FROM racing.race_results AS r ,
UNNEST(r.participants) AS p , UNNEST(p.splits) AS split_times WHERE p.name LIKE 'R%' GROUP BY
p.name ORDER BY total_race_time ASC;
```

Row name total_race_time

1 Rudisha 102.19999999999999

2 Rotich 103.6

Task 9. Lab question: Filtering within ARRAY values

You happened to see that the fastest lap time recorded for the 800 M race was 23.2 seconds, but you did not see which runner ran that particular lap. Create a query that returns that result.

- **Task:** Complete the partially written query:

```
#standardSQL SELECT p.name, split_time FROM racing.race_results AS r , r.participants AS p , p.splits AS
split_time WHERE split_time = ;
```

Possible solution:

```
#standardSQL SELECT p.name, split_time FROM racing.race_results AS r , UNNEST(r.participants) AS p ,
UNNEST(p.splits) AS split_time WHERE split_time = 23.2;
```

Row name split_time

1 Kipketer 23.2

Congratulations!

You've successfully ingested JSON datasets, created ARRAYS and STRUCTs, and unnested semi-structured data for insights.

Next steps / Learn more

For additional reading, refer to [Working with Arrays](#).

Copyright 2022 Google LLC All rights reserved. Google and the Google logo are trademarks of Google LLC. All other company and product names may be trademarks of the respective companies with which they are associated.